

**MPEI 2018-2019**

---

# **#22 – Bloom Filters (continuation)**

---

# Analysis

---

# Basic idea: Throwing Darts

- If we throw  $m$  darts into  $n$  equally likely targets, **what is the probability that a target gets at least one dart?**
- **In our case:**
  - **Targets** = bits/buckets
  - **Darts** = hash values of items

# Probabilities for one bit

- Initially all bits are set to zero.
- What is the probability of bit  $b_i = 1$  after using the first hash function when inserting an element ?
- Assuming hash function selects each array position with equal probability, is  $1/n$
- So, the probability of  $b_i = 0$  is  $1 - \frac{1}{n}$
- To insert the element this process is repeated  $k$  times
- Resulting in probability of having  $b_i = 0$  :

$$\left(1 - \frac{1}{n}\right)^k$$

# Probabilities for one bit (cont.)

- Since there are  $k$  hash functions and  $m$  elements to insert in the filter, the probability of having  $b_i = 0$  after handling all  $m$  elements (assuming independence) is:
  - $\left(1 - \frac{1}{n}\right)^{k m}$
  - $= \left[\left(1 - \frac{1}{n}\right)^m\right]^k = a^k$
- The probability of  $b_i = 1$  is  $1 - a^k$
- Related to the **probability that a target gets at least one dart**

# False positives and false negatives

Element in set ?	Result of membership test	Correct ?	Type of error
Yes	Yes (in set)	Yes	
	No	Error	False negative
No	Yes	Error	False positive
	No	Yes	

# Probability of a false positive

- There is a false positive error when we look up a bit string that is not in  $S$ , but find the corresponding  $k$  bits set to 1
- The probability of such event, after inserting  $m$  elements, is:

- $$p = \left[ 1 - \left( 1 - \frac{1}{n} \right)^{k m} \right]^k = \left( 1 - a^k \right)^k$$


- $$p \approx \left( 1 - e^{-km/n} \right)^k$$

- Using the approximation  $(1 - \epsilon)^{1/\epsilon} = 1/e$  (small  $\epsilon$ )



# Effect of k

- *Example:  $m = 1$  billion,  $n = 8$  billion*
  - $k = 1: (1 - e^{-1/8}) = 0.1175$
  - $k = 2: (1 - e^{-1/4})^2 = 0.0493$
- Adding a **second hash function** reduced false positives
- What happens as we keep increasing  $k$ ?



# “Optimal” value of $k$

- To determine the number of hash functions ( $k$ ) that minimizes the probability of error ( $p$ ) we minimize  $\log p$ , more tractable [ $\log = \ln$ ]
- $\log p = k \log (1 - a^k)$
- Finding derivative and setting to zero:
- $(1 - a^k) \log (1 - a^k) - a^k \log a^k = 0$
- With solution  $a^k = \frac{1}{2}$
- The best value of  $k$  is therefore
- $$k_{opt} = \frac{\log 1/2}{\log a} = \frac{\log 1/2}{m \log(1 - \frac{1}{n})} \approx \frac{n \log 2}{m} \approx \frac{0.693 n}{m}$$

# “Optimal” value of $k$

- As  $k_{opt}$  is not an integer, in practice we use the closest integer
- “Optimal” value of  $k$ :  $n/m \ln(2)$
- In our example ( $m = 1$  billion,  $n = 8$  billion)
  - Optimal  $k = 8 \ln(2) = 5.54 \approx 6$

# Lower bound for the error probability

- If we could have  $a^k = 1/2$  for an integer  $k_{opt}$ , the equation  $(1 - a^k)^k$  would lead to:
- $p_{opt} = \left(1 - \frac{1}{2}\right)^{k_{opt}}$
- $= \left(\frac{1}{2}\right)^{k_{opt}}$
- $= 2^{-k_{opt}}$
- That can be taken as the lower bound for the error probability (false positives)

# Value for n ?

- The required  $n$ , given  $m$  and a desired false positive probability  $p$  - and assuming the optimal value of  $k$  is used - can be computed by substituting the optimal value of  $k$  [  $=n/m \ln(2)$  ] in the probability expression

$$p \approx \left(1 - e^{-km/n}\right)^k$$

- $n = ?$ 
  - Homework

# One or more arrays ?

- Is it better to have **1** big **B** or **k** small **Bs**?
- $(1 - e^{-km/n})^k$  vs  $(1 - e^{-m/(n/k)})^k$
- It is the same
- But keeping **1** big **B** is simpler

# Perfect Hashing

---

- If the membership set is known in advance then better performance than with a standard Bloom Filter can be achieved using related techniques.
- For instance we can establish an arbitrarily low false positive rate by **using perfect hash functions** and limiting the size of the shared hash table.

# Bloom Filter: Wrap-up

---

- **Bloom filters guarantee no false negatives, and use limited memory**
  - Great for pre-processing before more expensive checks
- **Suitable for hardware implementation**
  - Hash function computations can be parallelized



# Trade-offs

---

- The error rate can be decreased
  - by increasing the number of hash functions
  - and the space allocated to store the table

# Final remarks

---

- Bloom Filters should be considered for programs where an imperfect set membership test could be helpfully applied to a large data set of unknown composition
- The great advantage of Bloom Filters over the use of single hash transforms is their speed and set error rate.
- Although the method can be applied to sets of any size, small sets are better dealt with by trees and heaps

# Counting Filters

---

# The origin

- A basic Bloom filter can only represent a set, but neither allows for querying the multiplicities of an item, nor does it support deleting entries
- The term *counting Bloom filter* is used to refer to **variants of Bloom filters** that represent **multisets** rather than sets
- Technically, a counting Bloom filter **extends a basic Bloom filter with width parameter  $w$  (bits)**.

# The basic idea

- To create a **Counting filter** the array positions are extended from a single bit to being an  $w$  bit counter.
- Counting filters provide a way to implement a *delete* operation on a Bloom filter without recreating the filter afresh
  - The original counting Bloom filter used cells with  $w=4$  only to support deletion, not to count elements
    - Making counting Bloom filters to use 3 to 4 times more space than basic Bloom filters.

# Changes to insert() and isMember()

---

- The **insert operation is extended** to *increment* the value of the buckets
- The lookup operation checks that each of the required buckets is non-zero.

# delete()

---

- New operation
- The delete operation consists of decrementing the value of each of the respective buckets
- **Deletions necessarily introduce false negative (FN) errors**
  - when you flip a set bit back to 0 that was part of a k bits from another item, the Bloom filter will no longer report that item as belonging to the set

# count()

---

- New operation
- Retrieving the count of an item of the set involves computing its set of counters and **returning the minimum value** as frequency estimate
- This query algorithm is also known as *minimum selection* (MS).



# Issues

- There exist two main issues with counting Bloom filters:
- Counter overflows
  - exists when the counter value reaches  $2w-1$  and cannot be incremented anymore.
  - One typically stops counting as opposed to overflowing and restarting at 0.
  - However, this strategy introduces *undercounts*, which we also refer to as FNs.
- The choice of  $w$ 
  - A large  $w$  quickly diminishes the space savings from using of a Bloom filter
    - There will also be a lot of unused space manifesting as unused zeros.
  - A small  $w$  may quickly lead to maximum counter values
  - Choosing the right value is a difficult trade-off that depends on the distribution of the data

# Limitations

---

- Counting filters have **limited scalability**
- Because the counting Bloom filter table cannot be expanded, the **maximal number of keys to be stored simultaneously in the filter must be known in advance.**
- Once the designed capacity of the table is exceeded, the false positive rate will grow rapidly as more keys are inserted

# Matlab

---

- Adapt insert and isMember
- Create the new functions:
  - delete()
  - count()
- Create a simple test

# Links used

---

- <http://matthias.vallentin.net/blog/2011/06/a-garden-variety-of-bloom-filters/>
- [https://en.wikipedia.org/wiki/Bloom\\_filter#Counting\\_filters](https://en.wikipedia.org/wiki/Bloom_filter#Counting_filters)

# Other techniques

---

- Variants of Bloom

- See, for example:

- <http://matthias.vallentin.net/course-work/cs270-s11.pdf>

- Cuckoo Hashing

- See:

- [http://www.lkozma.net/cuckoo hashing visualization/](http://www.lkozma.net/cuckoo_hashing_visualization/)

**Note to other teachers and users of these slides:** We would be delighted if you found this our material useful in giving your own lectures. Feel free to use these slides verbatim, or to modify them to fit your own needs. If you make use of a significant portion of these slides in your own lecture, please include this message, or a link to our web site: <http://www.mmds.org>

# Part of the slides adapted from: Mining Data Streams

---

Mining of Massive Datasets

Jure Leskovec, Anand Rajaraman, Jeff Ullman

Stanford University

<http://www.mmds.org>

